

loan-tap-logistic-regression

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report,
    roc_auc_score
from sklearn.metrics import precision_recall_curve, auc, roc_curve
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    f1_score
```

```
[ ]: from google.colab import files
uploaded = files.upload()
```

```
<IPython.core.display.HTML object>
Saving LoanTapData.csv to LoanTapData.csv
```

```
[ ]: df = pd.read_csv("LoanTapData.csv")
df.sample(5)
```

```
      loan_amnt      term  int_rate  installment  grade  sub_grade \
55376        25000  36 months     15.31      870.44    C       C2
176997       12000  60 months     17.86      303.81    D       D5
242595       20000  60 months     18.99      518.71    E       E1
330244       14125  36 months     18.75      515.99    D       D3
378456        7000  36 months     15.61      244.76    C       C5

                           emp_title  emp_length  home_ownership  annual_inc ... \
55376                  CSC        < 1 year      MORTGAGE   120000.0 ...
176997  Sales Representative    10+ years       RENT   120000.0 ...
242595                 carrier    10+ years      MORTGAGE   65000.0 ...
330244    Cox Communications      4 years      MORTGAGE  115000.0 ...
378456                 Manager    10+ years      MORTGAGE  109000.0 ...

      open_acc  pub_rec  revol_bal  revol_util  total_acc  initial_list_status \
55376       10        0     20359       57.8         27                  f
```

176997	9	0	7423	24.1	17	w
242595	10	0	25305	88.5	20	w
330244	13	0	4012	57.3	23	w
378456	11	0	8614	46.8	19	w
55376	INDIVIDUAL		1.0		0.0	
176997	INDIVIDUAL		1.0		0.0	
242595	INDIVIDUAL		3.0		0.0	
330244	INDIVIDUAL		4.0		0.0	
378456	INDIVIDUAL		2.0		0.0	
					address	
55376	263 Susan Stream Apt.	499\r\nKimberlytown, NM ...				
176997	PSC 8941, Box 5836\r\nAPO AP 00813					
242595	USCGC Cook\r\nFPO AP 70466					
330244	15435 Brown Mountains Apt.	016\r\nDanieltown, ...				
378456	67034 William Islands Apt.	134\r\nSouth Samant...				

[5 rows x 27 columns]

0.0.1 Data Exploration and Wrangling

```
[ ]: df.info()
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   loan_amnt        396030 non-null   int64  
 1   term              396030 non-null   object  
 2   int_rate          396030 non-null   float64 
 3   installment       396030 non-null   float64 
 4   grade             396030 non-null   object  
 5   sub_grade         396030 non-null   object  
 6   emp_title         373103 non-null   object  
 7   emp_length        377729 non-null   object  
 8   home_ownership    396030 non-null   object  
 9   annual_inc        396030 non-null   float64 
 10  verification_status 396030 non-null   object  
 11  issue_d           396030 non-null   object  
 12  loan_status        396030 non-null   object  
 13  purpose            396030 non-null   object  
 14  title              394274 non-null   object  
 15  dti                396030 non-null   float64
```

```

16 earliest_cr_line      396030 non-null object
17 open_acc              396030 non-null int64
18 pub_rec               396030 non-null int64
19 revol_bal             396030 non-null int64
20 revol_util            395754 non-null float64
21 total_acc              396030 non-null int64
22 initial_list_status   396030 non-null object
23 application_type       396030 non-null object
24 mort_acc               358235 non-null float64
25 pub_rec_bankruptcies  395495 non-null float64
26 address                396030 non-null object
dtypes: float64(7), int64(5), object(15)
memory usage: 81.6+ MB

```

```

[ ]:      loan_amnt      int_rate      installment      annual_inc \
count  396030.000000  396030.000000  396030.000000  3.960300e+05
mean    14113.888089   13.639400   431.849698  7.420318e+04
std     8357.441341    4.472157   250.727790  6.163762e+04
min     500.000000    5.320000   16.080000  0.000000e+00
25%    8000.000000   10.490000  250.330000  4.500000e+04
50%   12000.000000   13.330000  375.430000  6.400000e+04
75%   20000.000000   16.490000  567.300000  9.000000e+04
max    40000.000000   30.990000 1533.810000  8.706582e+06

          dti      open_acc      pub_rec      revol_bal \
count  396030.000000  396030.000000  396030.000000  3.960300e+05
mean    17.379514    11.311153   0.178191  1.584454e+04
std     18.019092    5.137649   0.530671  2.059184e+04
min     0.000000    0.000000   0.000000  0.000000e+00
25%    11.280000    8.000000   0.000000  6.025000e+03
50%   16.910000   10.000000   0.000000  1.118100e+04
75%   22.980000   14.000000   0.000000  1.962000e+04
max    9999.000000   90.000000  86.000000  1.743266e+06

      revol_util      total_acc      mort_acc  pub_rec_bankruptcies
count  395754.000000  396030.000000  358235.000000  395495.000000
mean    53.791749    25.414744   1.813991   0.121648
std     24.452193    11.886991   2.147930   0.356174
min     0.000000    2.000000   0.000000   0.000000
25%    35.800000    17.000000   0.000000   0.000000
50%    54.800000    24.000000   1.000000   0.000000
75%    72.900000    32.000000   3.000000   0.000000
max    892.300000   151.000000  34.000000   8.000000

```

```

[ ]: # No. of Rows and Columns
print("No. of Rows:", df.shape[0])
print("No. of Columns:", df.shape[1])

```

```
No. of Rows: 396030
```

```
No. of Columns: 27
```

```
[ ]: df.isna().sum().sum()
# we can drop the null values, but we are not going to drop it, instead we are ↴going to impute it
```

```
[ ]: 81590
```

```
[ ]: # Distribution of Target labels
df["loan_status"].value_counts(normalize=True)
```

```
[ ]: loan_status
Fully Paid      0.803871
Charged Off    0.196129
Name: proportion, dtype: float64
```

- There is a good amount of imbalance in the target outcomes, this implies accuracy metric will not be a good option to evaluate our model.
- First let I will fit a model with the same imbalanced data and check for evaluation metrics then after that I will fit model on the dataset to which I am going to oversample it for reducing the imbalance of the minority class.

```
[ ]: # correlation Matrix
df.select_dtypes(include="number").corr(method="spearman")
```

```
[ ]:          loan_amnt  int_rate  installment  annual_inc       dti \
loan_amnt        1.000000  0.131432    0.968334   0.488566  0.053118
int_rate         0.131432  1.000000    0.137293  -0.096648  0.172123
installment      0.968334  0.137293    1.000000   0.470464  0.055538
annual_inc       0.488566 -0.096648    0.470464   1.000000 -0.202856
dti              0.053118  0.172123    0.055538  -0.202856  1.000000
open_acc         0.215244  0.004181    0.207828   0.240468  0.322745
pub_rec          -0.100435  0.072204   -0.093357  -0.046322 -0.042220
revol_bal        0.469646  0.005948    0.460147   0.393232  0.249720
revol_util       0.104708  0.303990    0.131628   0.060172  0.185130
total_acc        0.237115 -0.050880    0.216667   0.334453  0.236925
mort_acc         0.231240 -0.102962    0.201861   0.378772 -0.048033
pub_rec_bankruptcies -0.108661  0.060848   -0.102922  -0.071770 -0.032790
```

```
[ ]:          open_acc  pub_rec  revol_bal  revol_util  total_acc \
loan_amnt        0.215244 -0.100435  0.469646   0.104708  0.237115
int_rate         0.004181  0.072204  0.005948   0.303990 -0.050880
installment      0.207828 -0.093357  0.460147   0.131628  0.216667
annual_inc       0.240468 -0.046322  0.393232   0.060172  0.334453
dti              0.322745 -0.042220  0.249720   0.185130  0.236925
open_acc         1.000000 -0.019111  0.364325   -0.139233  0.672232
pub_rec          -0.019111  1.000000 -0.209249  -0.095391  0.033317
```

revol_bal	0.364325	-0.209249	1.000000	0.419506	0.294369
revol_util	-0.139233	-0.095391	0.419506	1.000000	-0.104718
total_acc	0.672232	0.033317	0.294369	-0.104718	1.000000
mort_acc	0.142336	0.031729	0.239253	0.008440	0.404738
pub_rec_bankruptcies	-0.025244	0.862245	-0.205400	-0.091208	0.041582
		mort_acc	pub_rec_bankruptcies		
loan_amnt	0.231240		-0.108661		
int_rate	-0.102962		0.060848		
installment	0.201861		-0.102922		
annual_inc	0.378772		-0.071770		
dti	-0.048033		-0.032790		
open_acc	0.142336		-0.025244		
pub_rec	0.031729		0.862245		
revol_bal	0.239253		-0.205400		
revol_util	0.008440		-0.091208		
total_acc	0.404738		0.041582		
mort_acc	1.000000		0.039623		
pub_rec_bankruptcies	0.039623		1.000000		

```
[ ]: df.select_dtypes(include="number").corr(method="pearson")
```

	loan_amnt	int_rate	installment	annual_inc	dti	\
loan_amnt	1.000000	0.168921	0.953929	0.336887	0.016636	
int_rate	0.168921	1.000000	0.162758	-0.056771	0.079038	
installment	0.953929	0.162758	1.000000	0.330381	0.015786	
annual_inc	0.336887	-0.056771	0.330381	1.000000	-0.081685	
dti	0.016636	0.079038	0.015786	-0.081685	1.000000	
open_acc	0.198556	0.011649	0.188973	0.136150	0.136181	
pub_rec	-0.077779	0.060986	-0.067892	-0.013720	-0.017639	
revol_bal	0.328320	-0.011280	0.316455	0.299773	0.063571	
revol_util	0.099911	0.293659	0.123915	0.027871	0.088375	
total_acc	0.223886	-0.036404	0.202430	0.193023	0.102128	
mort_acc	0.222315	-0.082583	0.193694	0.236320	-0.025439	
pub_rec_bankruptcies	-0.106539	0.057450	-0.098628	-0.050162	-0.014558	
	open_acc	pub_rec	revol_bal	revol_util	total_acc	\
loan_amnt	0.198556	-0.077779	0.328320	0.099911	0.223886	
int_rate	0.011649	0.060986	-0.011280	0.293659	-0.036404	
installment	0.188973	-0.067892	0.316455	0.123915	0.202430	
annual_inc	0.136150	-0.013720	0.299773	0.027871	0.193023	
dti	0.136181	-0.017639	0.063571	0.088375	0.102128	
open_acc	1.000000	-0.018392	0.221192	-0.131420	0.680728	
pub_rec	-0.018392	1.000000	-0.101664	-0.075910	0.019723	
revol_bal	0.221192	-0.101664	1.000000	0.226346	0.191616	
revol_util	-0.131420	-0.075910	0.226346	1.000000	-0.104273	
total_acc	0.680728	0.019723	0.191616	-0.104273	1.000000	

```

mort_acc          0.109205  0.011552  0.194925  0.007514  0.381072
pub_rec_bankruptcies -0.027732  0.699408 -0.124532 -0.086751  0.042035

mort_acc  pub_rec_bankruptcies
loan_amnt      0.222315      -0.106539
int_rate        -0.082583      0.057450
installment     0.193694      -0.098628
annual_inc       0.236320      -0.050162
dti             -0.025439      -0.014558
open_acc         0.109205      -0.027732
pub_rec          0.011552      0.699408
revol_bal        0.194925      -0.124532
revol_util       0.007514      -0.086751
total_acc        0.381072      0.042035
mort_acc         1.000000      0.027239
pub_rec_bankruptcies  0.027239      1.000000

```

- As you can see both spearman and pearson correlation values are almost equal which says that there might be a presence of slight to no non-linearity in the features.
- The installment amount and loan amount are highly correlated which has correlation of over 0.95 so we can remove any one of the 2 features.

```
[ ]: # Dropping installemnt column
df.drop(columns=["installment"], inplace=True)

[ ]: df["term"].unique()

[ ]: array([' 36 months', ' 60 months'], dtype=object)

[ ]: # Cleaning individual columns
# term
df["term"] = df["term"].str.strip()
map = {"36 months": 36, "60 months": 60}
df["term"] = df["term"].map(map)

[ ]: # grade
df["grade"].value_counts()

[ ]: grade
B    116018
C    105987
A    64187
D    63524
E    31488
F    11772
G    3054
Name: count, dtype: int64
```

```
[ ]: df["emp_length"].value_counts()
map = {"10+ years":10,"2 years":2,"< 1 year":1,"3 years":3,"5 years":5,"6 years":6,
       "4 years":4,"1 year":1,"7 years":7,"8 years":8,"9 years":9}
df["emp_length"] = df["emp_length"].map(map)
```

```
[ ]: df["emp_length"].value_counts()
```

```
[ ]: emp_length
10.0      126041
1.0       57607
2.0       35827
3.0       31665
5.0       26495
4.0       23952
6.0       20841
7.0       20819
8.0       19168
9.0       15314
Name: count, dtype: int64
```

```
[ ]: df["home_ownership"].value_counts()
```

```
[ ]: home_ownership
MORTGAGE     198348
RENT         159790
OWN          37746
OTHER         112
NONE          31
ANY            3
Name: count, dtype: int64
```

```
[ ]: df.loc[(df["home_ownership"]=="NONE") | (df["home_ownership"]=="ANY"), "home_ownership"] = "OTHER"
```

```
[ ]: df["home_ownership"].value_counts()
```

```
[ ]: home_ownership
MORTGAGE     198348
RENT         159790
OWN          37746
OTHER         146
Name: count, dtype: int64
```

```
[ ]: df["verification_status"].value_counts()
```

```
[ ]: verification_status  
Verified          139563  
Source Verified    131385  
Not Verified       125082  
Name: count, dtype: int64
```

```
[ ]: # since the full address is not of use  
df["zipcode"] = df["address"].apply(lambda x:x[-5:])
```

```
[ ]: # We extracted zip code from address column, now address is of no use, that is  
# why we are dropping it  
df.drop(columns=["address"], inplace=True)
```

0.0.2 Converting Data Types of columns

```
[ ]: df["grade"] = df["grade"].astype("category")  
df["sub_grade"] = df["sub_grade"].astype("category")  
df["verification_status"] = df["verification_status"].astype("category")
```

```
[ ]: df["issue_d"] = pd.to_datetime(df["issue_d"], format='%b-%y')  
df["earliest_cr_line"] = pd.to_datetime(df["earliest_cr_line"], format='%b-%y')
```

```
[ ]: df["title"].value_counts()
```

```
[ ]: title  
Debt consolidation          152472  
Credit card refinancing     51487  
Home improvement            15264  
Other                        12930  
Debt Consolidation          11608  
...  
buisness startup              1  
Medical-Lap Band             1  
Debt consolidation to improve my life! 1  
one debt only                 1  
Toxic Debt Payoff             1  
Name: count, Length: 48787, dtype: int64
```

```
[ ]: df["purpose"].value_counts()
```

```
[ ]: purpose  
debt_consolidation      234507  
credit_card                83019  
home_improvement         24030  
other                      21185  
major_purchase              8790  
small_business               5701
```

```
car                    4697
medical                4196
moving                  2854
vacation                2452
house                   2201
wedding                 1812
renewable_energy         329
educational              257
Name: count, dtype: int64
```

```
[ ]: # As you can see Title and purpose column are giving same information, so we can drop one of that
```

```
[ ]: df.drop(columns=["title"], inplace=True)
```

```
[ ]: # Initial list status
df["initial_list_status"].value_counts()
df["initial_list_status"].astype("category")
```

```
[ ]: 0          w
1          f
2          f
3          f
4          f
..
396025    w
396026    f
396027    f
396028    f
396029    f
Name: initial_list_status, Length: 396030, dtype: category
Categories (2, object): ['f', 'w']
```

```
[ ]: df["application_type"].value_counts()
df["application_type"].astype("category")
```

```
[ ]: 0          INDIVIDUAL
1          INDIVIDUAL
2          INDIVIDUAL
3          INDIVIDUAL
4          INDIVIDUAL
..
396025    INDIVIDUAL
396026    INDIVIDUAL
396027    INDIVIDUAL
396028    INDIVIDUAL
396029    INDIVIDUAL
```

```
Name: application_type, Length: 396030, dtype: category  
Categories (3, object): ['DIRECT_PAY', 'INDIVIDUAL', 'JOINT']
```

```
[ ]: # By converting columns into categorical I have reduced the memory usage from  
     ↪ 82 Mb to 68 Mb  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 396030 entries, 0 to 396029  
Data columns (total 25 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   loan_amnt        396030 non-null  int64    
 1   term             396030 non-null  int64    
 2   int_rate          396030 non-null  float64  
 3   grade            396030 non-null  category  
 4   sub_grade         396030 non-null  category  
 5   emp_title         373103 non-null  object    
 6   emp_length        377729 non-null  float64  
 7   home_ownership    396030 non-null  object    
 8   annual_inc        396030 non-null  float64  
 9   verification_status 396030 non-null  category  
 10  issue_d           396030 non-null  datetime64[ns]  
 11  loan_status       396030 non-null  object    
 12  purpose           396030 non-null  object    
 13  dti               396030 non-null  float64  
 14  earliest_cr_line  396030 non-null  datetime64[ns]  
 15  open_acc          396030 non-null  int64    
 16  pub_rec            396030 non-null  int64    
 17  revol_bal          396030 non-null  int64    
 18  revol_util         395754 non-null  float64  
 19  total_acc          396030 non-null  int64    
 20  initial_list_status 396030 non-null  object    
 21  application_type   396030 non-null  object    
 22  mort_acc           358235 non-null  float64  
 23  pub_rec_bankruptcies 395495 non-null  float64  
 24  zipcode            396030 non-null  object    
dtypes: category(3), datetime64[ns](2), float64(7), int64(6), object(7)  
memory usage: 67.6+ MB
```

```
[ ]: df["issue_year"] = df["issue_d"].dt.year  
df["issue_month"] = df["issue_d"].dt.month  
df["earliest_cr_year"] = df["earliest_cr_line"].dt.year  
df["earliest_cr_month"] = df["earliest_cr_line"].dt.month
```

```
[ ]: df.drop(columns=["issue_d", "earliest_cr_line"], axis=1, inplace=True)
```

Handling Outlier and Null values

```
[ ]: df_copy = df.copy()

[ ]: num_cols = ["loan_amnt","int_rate","annual_inc","dti","revol_bal","revol_util",
               "total_acc","pub_rec","open_acc","mort_acc","pub_rec_bankruptcies"]

def remove_outliers(df, cols):
    for col in cols:
        q1 = df[col].quantile(0.25)
        q3 = df[col].quantile(0.75)
        IQR = q3-q1
        lower_bound = q1-(1.5*IQR)
        upper_bound = q3+(1.5*IQR)
        df = df[(df[col]>=lower_bound) & (df[col]<=upper_bound)]
    return df

[ ]: df_copy = remove_outliers(df_copy, num_cols)

[ ]: # After removing outliers, There is a lot of information loss, almost 36% data loss,
     # which is very significant
# But still we have to remove outliers, but now we discard less important columns from the list
df_copy.shape

[ ]: (253377, 27)

[ ]: # creating one more copy of DataFrame
df_new = df.copy()

[ ]: num_cols = ["loan_amnt","int_rate","annual_inc","dti","revol_bal","revol_util"]

def remove_outliers(df, cols):
    for col in cols:
        q1 = df[col].quantile(0.25)
        q3 = df[col].quantile(0.75)
        IQR = q3-q1
        lower_bound = q1-(1.5*IQR)
        upper_bound = q3+(1.5*IQR)
        df = df[(df[col]>=lower_bound) & (df[col]<=upper_bound)]
    return df

[ ]: df_new = remove_outliers(df_new, num_cols)

[ ]: df_new.isnull().sum()
```

```
[ ]: loan_amnt          0
      term             0
      int_rate          0
      grade            0
      sub_grade         0
      emp_title         21099
      emp_length        17289
      home_ownership    0
      annual_inc        0
      verification_status 0
      loan_status        0
      purpose            0
      dti                0
      open_acc           0
      pub_rec            0
      revol_bal          0
      revol_util         0
      total_acc          0
      initial_list_status 0
      application_type   0
      mort_acc           34714
      pub_rec_bankruptcies 483
      zipcode            0
      issue_year          0
      issue_month         0
      earliest_cr_year    0
      earliest_cr_month   0
      dtype: int64
```

```
[ ]: df_new["emp_title"].nunique()
```

```
[ ]: 159668
```

- Since null values in pub_rec_bankruptcies are very less we can just drop those
- we have to impute null values in mort_acc and emp_length
- emp_title column has so many distinct values compared to its null values so that we can remove this column for now for our analysis

```
[ ]: df_new.groupby(["emp_title","loan_status"]).size().unstack(fill_value=0).
    ↪sort_values(by="Fully Paid",ascending=False)
```

	Charged Off	Fully Paid
loan_status		
emp_title		
Teacher	798	3344
Manager	818	2954
Registered Nurse	342	1376
Supervisor	361	1335

RN	341	1322
...
City of Northport	1	0
City of Northport Northport, Al	1	0
Warehouse Laborer	1	0
Warehouse Labor	1	0
Doumak Inc	1	0

[159668 rows x 2 columns]

- The above dataframe shows that teacher, manager and Nurse are the most common professions that had been afforded the loan

```
[ ]: df_new.groupby(["purpose", "loan_status"]).size().unstack().
    ↪sort_values(by="Fully Paid", ascending=False)
```

	Charged Off	Fully Paid
purpose		
debt_consolidation	44077	168642
credit_card	12521	61511
home_improvement	3552	17079
other	4033	15166
major_purchase	1324	6644
car	579	3765
small_business	1365	3405
medical	812	2983
moving	607	2008
vacation	436	1853
wedding	203	1511
house	359	1500
renewable_energy	70	233
educational	38	200

- Most loans were given out for debt_consolidation, credit card refinancing.

```
[ ]: df_new.isna().sum()
```

loan_amnt	0
term	0
int_rate	0
grade	0
sub_grade	0
emp_title	21099
emp_length	17289
home_ownership	0
annual_inc	0
verification_status	0
loan_status	0

```
purpose          0
dti             0
open_acc        0
pub_rec         0
revol_bal       0
revol_util      0
total_acc       0
initial_list_status 0
application_type 0
mort_acc        34714
pub_rec_bankruptcies 483
zipcode         0
issue_year      0
issue_month     0
earliest_cr_year 0
earliest_cr_month 0
dtype: int64
```

```
[ ]: # we are removing pub_rec_bankruptcies because it has less number of null values
df_new.dropna(subset=["pub_rec_bankruptcies"], inplace=True)
```

```
[ ]: # we are also removing emp_title because of large number of unique values and
    ↪it is less important column
df_new.drop(columns=["emp_title"], inplace=True)
```

```
[ ]: df_new.isna().sum()
```

```
loan_amnt        0
term            0
int_rate         0
grade           0
sub_grade        0
emp_length       17289
home_ownership   0
annual_inc       0
verification_status 0
loan_status       0
purpose          0
dti              0
open_acc          0
pub_rec          0
revol_bal         0
revol_util        0
total_acc         0
initial_list_status 0
application_type 0
mort_acc        34231
```

```

pub_rec_bankruptcies      0
zipcode                   0
issue_year                0
issue_month               0
earliest_cr_year          0
earliest_cr_month         0
dtype: int64

```

Handling mort_acc and emp_length columns, below strategies can be used for that:

- For emp_length we can use median to fill it or even mean.
- For mort_acc we can fill it with mean, but not the mean of whole data, we can group it by the total accounts and then take mean of mort_acc since if total accounts are more then there is more probability that the person will have more mort_acc.

```
[ ]: df_new["emp_length"].median()
```

```
[ ]: 6.0
```

```
[ ]: df_new["emp_length"] = df_new["emp_length"].fillna(df_new["emp_length"].
   ↪median())
```

```
[ ]: # Now imputing mort_acc with the mean of mort_acc per total_acc
df_mort=round(df_new.groupby('total_acc')[['mort_acc']].mean(),0).reset_index()
```

```
[ ]: df_mort
```

	total_acc	mort_acc
0	2	0.0
1	3	0.0
2	4	0.0
3	5	0.0
4	6	0.0
..
107	116	1.0
108	117	0.0
109	118	1.0
110	124	1.0
111	135	3.0

```
[112 rows x 2 columns]
```

```
[ ]: # creating copy
df_new1=df_new.copy()
```

```
[ ]: df_merged = df_new.merge(df_mort, how="left", on="total_acc",
   ↪suffixes=( "", "_mean"))
```

```
[ ]: df_merged['mort_acc'].fillna(df_merged['mort_acc_mean']).isna().value_counts()
```

```
[ ]: mort_acc
```

```
False    355993
```

```
Name: count, dtype: int64
```

```
[ ]: df_merged['mort_acc'] = df_merged['mort_acc'].fillna(df_merged['mort_acc_mean'])
```

```
[ ]: df_merged.drop(columns=['mort_acc_mean'], inplace=True)
```

```
[ ]: df_merged
```

```
[ ]:      loan_amnt  term  int_rate grade sub_grade  emp_length home_ownership \
0        10000    36    11.44     B       B4         10.0        RENT
1        8000     36    11.99     B       B5          4.0      MORTGAGE
2       15600    36    10.49     B       B3          1.0        RENT
3        7200    36     6.49     A       A2          6.0        RENT
4       24375    60    17.27     C       C5          9.0      MORTGAGE
...
355988     6000    36    13.11     B       B4          5.0        RENT
355989    10000    60    10.99     B       B4          2.0        RENT
355990     5000    36     9.99     B       B1         10.0        RENT
355991    21000    60    15.31     C       C2         10.0      MORTGAGE
355992     2000    36    13.61     C       C2         10.0        RENT

      annual_inc verification_status  loan_status ... total_acc \
0      117000.0        Not Verified  Fully Paid ...      25
1      65000.0        Not Verified  Fully Paid ...      27
2      43057.0      Source Verified  Fully Paid ...      26
3      54000.0        Not Verified  Fully Paid ...      13
4      55000.0           Verified  Charged Off ...      43
...
355988     64000.0        Not Verified  Fully Paid ...       9
355989     40000.0      Source Verified  Fully Paid ...      23
355990     56500.0           Verified  Fully Paid ...      23
355991     64000.0           Verified  Fully Paid ...      20
355992     42996.0           Verified  Fully Paid ...      19

      initial_list_status application_type  mort_acc  pub_rec_bankruptcies \
0                  w      INDIVIDUAL      0.0            0.0
1                  f      INDIVIDUAL      3.0            0.0
2                  f      INDIVIDUAL      0.0            0.0
3                  f      INDIVIDUAL      0.0            0.0
4                  f      INDIVIDUAL      1.0            0.0
...
355988      ...                  w      INDIVIDUAL      0.0            ...
355989      ...                  w      INDIVIDUAL      0.0            0.0
```

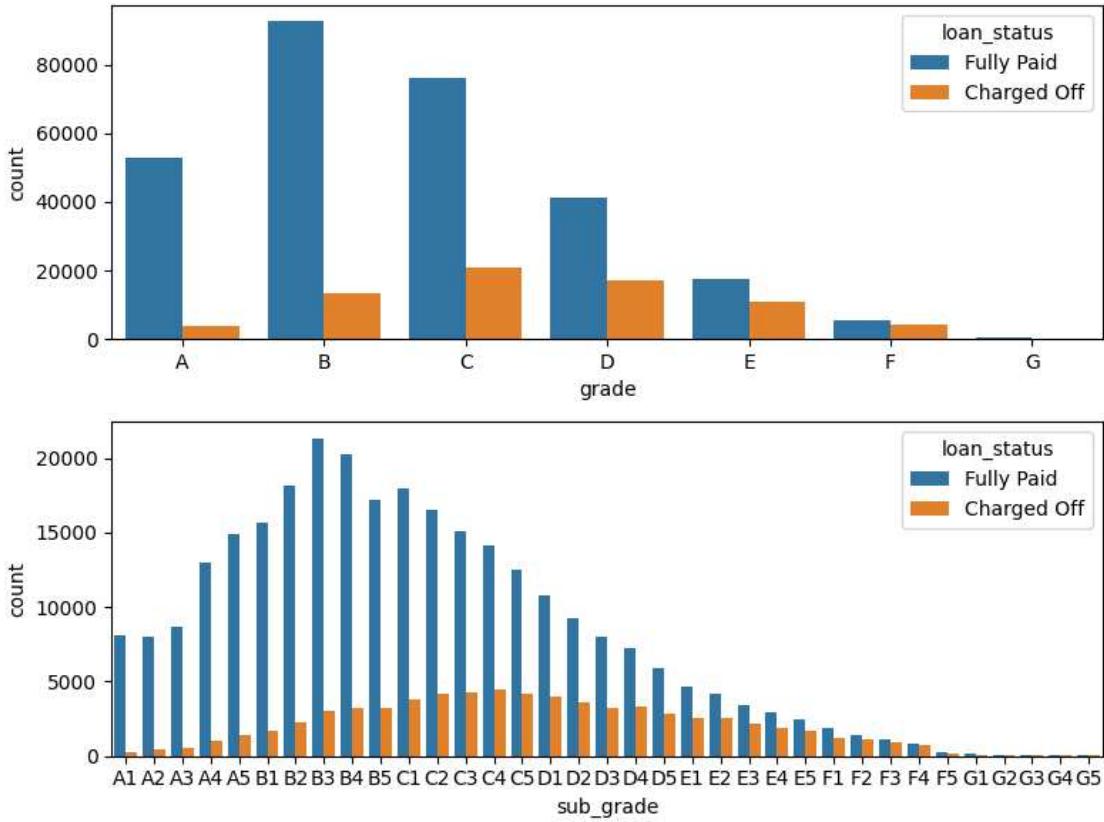
355990	f	INDIVIDUAL	0.0	0.0
355991	f	INDIVIDUAL	5.0	0.0
355992	f	INDIVIDUAL	1.0	0.0
zipcode	issue_year	issue_month	earliest_cr_year	earliest_cr_month
0	22690	2015	1	1990
1	05113	2015	1	2004
2	05113	2015	1	2007
3	00813	2014	11	2006
4	11650	2013	4	1999
...
355988	05113	2013	3	1991
355989	30723	2015	10	2004
355990	70466	2013	10	1997
355991	29597	2012	8	1990
355992	48052	2010	6	1998

[355993 rows x 26 columns]

Visualiazation and Graphic Analysis

Count plots of grade and sub-grade

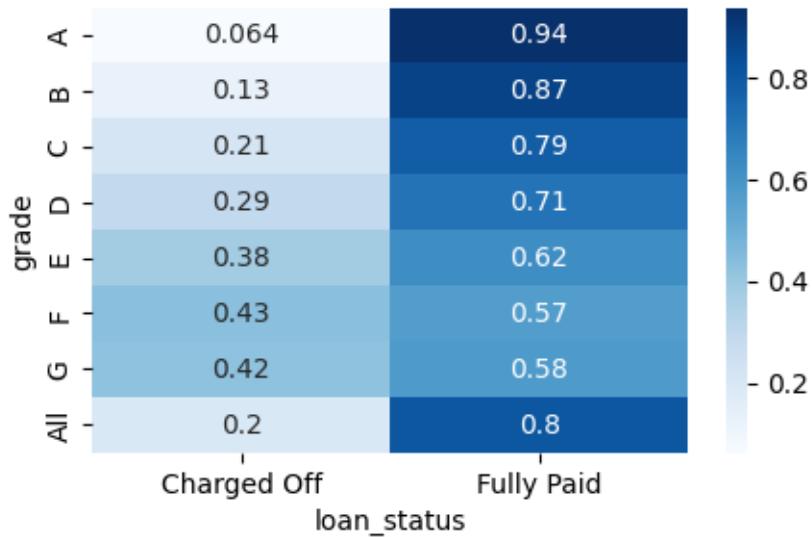
```
[ ]: fig,axes=plt.subplots(2,1,figsize=(8,6))
sns.countplot(data=df_merged,x='grade',hue='loan_status',ax=axes[0])
sns.countplot(data=df_merged,x='sub_grade',hue='loan_status',ax=axes[1])
plt.tight_layout()
plt.show()
```



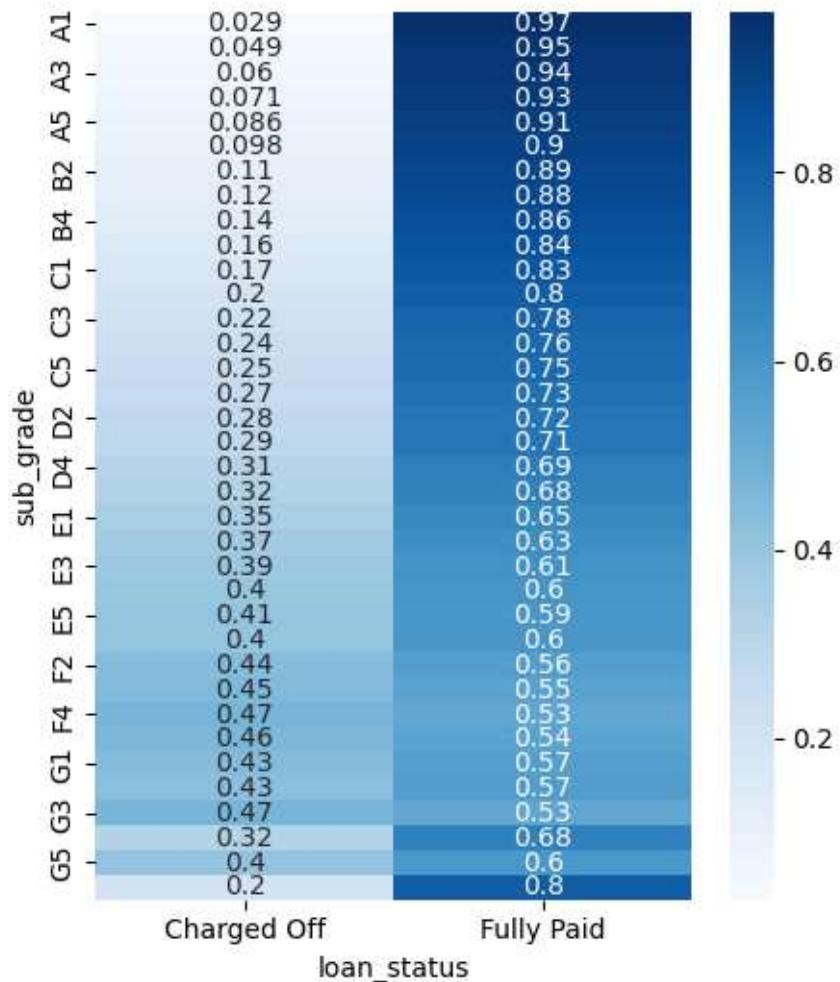
Observations from the above countplots:

- Most number of people in grade B paid off the loan.
- In Grade B most people in sub_grade B3 have paid off their loans.
- Very less number of people are categorised in G grade.
- The percentage of defaulters in sub_grade A1 is the least, implying genuine people.

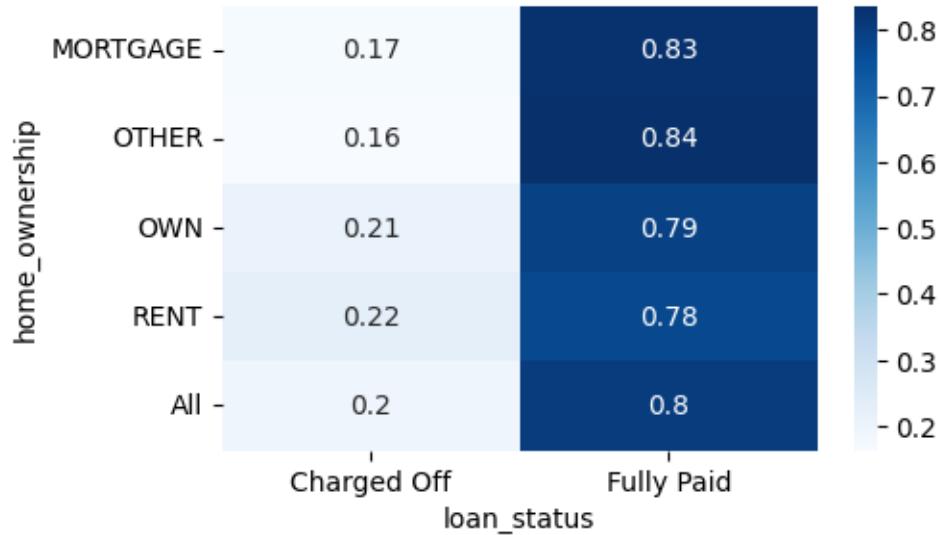
```
[ ]: plt.figure(figsize=(5,3))
sns.heatmap(pd.crosstab(df_merged['grade'],df_merged['loan_status'],normalize='index',margins=True),cmap='Blues')
plt.show()
```



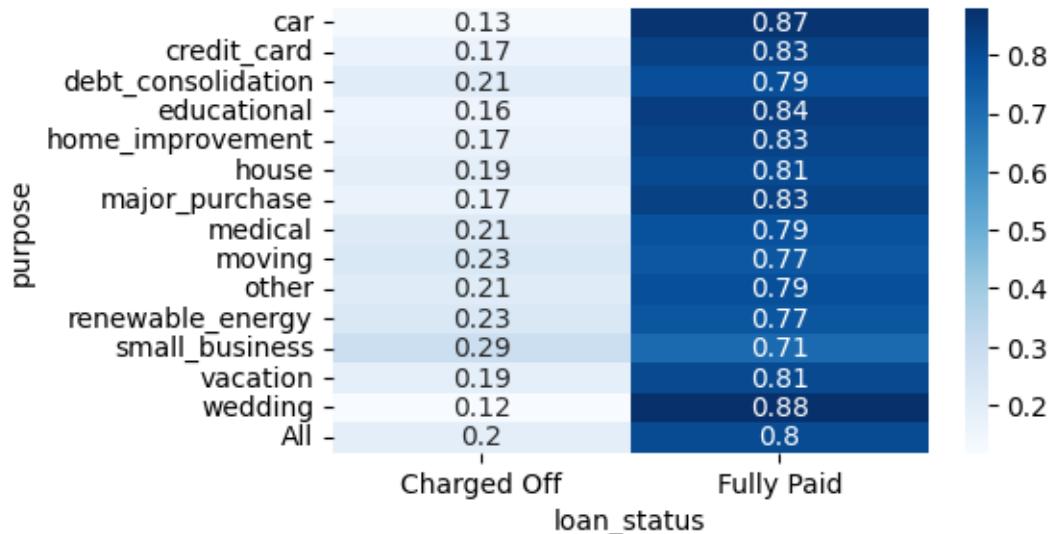
```
[ ]: plt.figure(figsize=(5,6))
sns.heatmap(pd.
    ↪crosstab(df_merged['sub_grade'],df_merged['loan_status'],normalize='index',margins=True),
    cmap='Blues')
plt.show()
```



```
[ ]: plt.figure(figsize=(5,3))
sns.heatmap(pd.
    crosstab(df_merged['home_ownership'],df_merged['loan_status'],normalize='index',margins=True,
    cmap='Blues')
plt.show()
```



```
[ ]: plt.figure(figsize=(5,3))
sns.heatmap(pd.
    ↪crosstab(df_merged['purpose'],df_merged['loan_status'],normalize='index',margins=True),cmap='Blues')
plt.show()
```



Observations from the above heatmaps:

- 94% of people marked in grade A paid back the loan.
- The customers who have opted for joint loans are more likely to pay back.

- The people who have taken loans for Wedding, Car and educational purposes are most likely to pay back the loan.
- People who have taken loan for small business are least likely to pay back the loan.

Feature Engineering

[]: df_merged

```
[ ]:      loan_amnt  term  int_rate grade sub_grade  emp_length home_ownership \
0        10000    36    11.44     B      B4         10.0       RENT
1         8000    36    11.99     B      B5          4.0      MORTGAGE
2        15600    36    10.49     B      B3          1.0       RENT
3         7200    36     6.49     A      A2          6.0       RENT
4        24375    60    17.27     C      C5          9.0      MORTGAGE
...
355988     6000    36    13.11     B      B4          5.0       RENT
355989    10000    60    10.99     B      B4          2.0       RENT
355990     5000    36     9.99     B      B1         10.0       RENT
355991    21000    60    15.31     C      C2         10.0      MORTGAGE
355992     2000    36    13.61     C      C2         10.0       RENT

      annual_inc verification_status  loan_status ... total_acc \
0      117000.0      Not Verified  Fully Paid ...      25
1      65000.0      Not Verified  Fully Paid ...      27
2      43057.0      Source Verified  Fully Paid ...      26
3      54000.0      Not Verified  Fully Paid ...      13
4      55000.0      Verified     Charged Off ...      43
...
355988    64000.0      Not Verified  Fully Paid ...       9
355989    40000.0      Source Verified  Fully Paid ...      23
355990    56500.0      Verified     Fully Paid ...      23
355991    64000.0      Verified     Fully Paid ...      20
355992    42996.0      Verified     Fully Paid ...      19

      initial_list_status application_type  mort_acc  pub_rec_bankruptcies \
0                  w      INDIVIDUAL      0.0           0.0
1                  f      INDIVIDUAL      3.0           0.0
2                  f      INDIVIDUAL      0.0           0.0
3                  f      INDIVIDUAL      0.0           0.0
4                  f      INDIVIDUAL      1.0           0.0
...
355988     ...          ...          ...          ...
355989     ...          ...          ...          ...
355990     ...          ...          ...          ...
355991     ...          ...          ...          ...
355992     ...          ...          ...          ...

      zipcode  issue_year issue_month earliest_cr_year  earliest_cr_month
```

```

0      22690    2015      1      1990      6
1      05113    2015      1      2004      7
2      05113    2015      1      2007      8
3      00813    2014     11      2006      9
4      11650    2013      4      1999      3
...
355988  05113    2013      3      1991     11
355989  30723    2015     10      2004     11
355990  70466    2013     10      1997      3
355991  29597    2012      8      1990     11
355992  48052    2010      6      1998      9

```

[355993 rows x 26 columns]

```

[ ]: # Loan status
map={'Charged Off':1,'Fully Paid':0}
df_merged['loan_status']=df_merged['loan_status'].map(map)

[ ]: # Initial_list_status
map={'w':0,'f':1}
df_merged['initial_list_status']=df_merged['initial_list_status'].map(map)

[ ]: x = df_merged.drop(columns=['loan_status'])
y = df_merged['loan_status']

[ ]: x_train_cv, x_test, y_train_cv, y_test = train_test_split(x, y, test_size=0.25,random_state=1)
x_train, x_test_cv, y_train, y_test_cv = train_test_split(x_train_cv,y_train_cv, test_size=0.25, random_state=1)

[ ]: # Target encoding of categorical columns (I have included zip code,
# year and month as they also have limited categorical values and are notcontinuous with more number of values)
encoding_cols = ['grade','sub_grade','home_ownership','verification_status','purpose',
                 'application_type','zipcode']
# Target encoding should only be applied to category or object dtypes, it won'twork on continuous data
# type like int or float.
# In our dataset zipcode is an object and not int or float.

[ ]: !pip install category_encoders
from category_encoders import TargetEncoder

encoder=TargetEncoder()

for i in encoding_cols:

```

```

x_train[i]=encoder.fit_transform(x_train[i],y_train)
x_test_cv[i]=encoder.transform(x_test_cv[i])
x_test[i]=encoder.transform(x_test[i])

```

```

Requirement already satisfied: category_encoders in
/usr/local/lib/python3.10/dist-packages (2.6.3)
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.10/dist-
packages (from category_encoders) (1.25.2)
Requirement already satisfied: scikit-learn>=0.20.0 in
/usr/local/lib/python3.10/dist-packages (from category_encoders) (1.2.2)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.10/dist-
packages (from category_encoders) (1.11.4)
Requirement already satisfied: statsmodels>=0.9.0 in
/usr/local/lib/python3.10/dist-packages (from category_encoders) (0.14.2)
Requirement already satisfied: pandas>=1.0.5 in /usr/local/lib/python3.10/dist-
packages (from category_encoders) (2.0.3)
Requirement already satisfied: patsy>=0.5.1 in /usr/local/lib/python3.10/dist-
packages (from category_encoders) (0.5.6)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.5->category_encoders)
(2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas>=1.0.5->category_encoders) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-
packages (from pandas>=1.0.5->category_encoders) (2024.1)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages
(from patsy>=0.5.1->category_encoders) (1.16.0)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-
packages (from scikit-learn>=0.20.0->category_encoders) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-
learn>=0.20.0->category_encoders) (3.5.0)
Requirement already satisfied: packaging>=21.3 in
/usr/local/lib/python3.10/dist-packages (from
statsmodels>=0.9.0->category_encoders) (24.1)

```

```

[ ]: # MinMax scaler for numerical columns
from sklearn.preprocessing import MinMaxScaler
numeric_cols = [
    'loan_amnt', 'term', 'int_rate', 'emp_length', 'annual_inc', 'dti', 'open_acc',
    'pub_rec', 'revol_bal', 'revol_util', 'total_acc', 'mort_acc', 'pub_rec_bankruptcies',
    'issue_year', 'issue_month', 'earliest_cr_year', 'earliest_cr_month']
minmax = MinMaxScaler()

```

```
[ ]: x_train_scaled = minmax.fit_transform(x_train[numeric_cols])
x_test_cv_scaled=minmax.transform(x_test_cv[numeric_cols])
x_test_scaled=minmax.transform(x_test[numeric_cols])

[ ]: x_train[numeric_cols]=x_train_scaled
x_test_cv[numeric_cols]=x_test_cv_scaled
x_test[numeric_cols]=x_test_scaled

[ ]: table_styles = {
    'cerulean_palette': [
        {'selector': 'th',
         'props': [('background-color', '#4E79A7'),
                    ('color', 'white'),
                    ('font-family', 'verdana')]}
    ], [
        {'selector': 'td',
         'props': [('background-color', '#A0CBE8'),
                    ('color', '#000000'),
                    ('font-family', 'verdana')]}
    ]
}
display(x_train.head(10).style.
        ↪set_table_styles(table_styles['cerulean_palette']).set_caption("Scaled Final
        ↪Data"))
```

<pandas.io.formats.style.Styler at 0x7b8ba7d4bcd0>

Modeling

```
[ ]: # Let's fit the logistic Regression model and check

model = LogisticRegression(random_state=1,max_iter=500,n_jobs=-1)
model.fit(x_train,y_train)

[ ]: LogisticRegression(max_iter=500, n_jobs=-1, random_state=1)

[ ]: y_pred=model.predict(x_test)
print(f" Accuracy score of the {model} = {round(model.
        ↪score(x_test,y_test)*100,2)}")
```

Accuracy score of the LogisticRegression(max_iter=500, n_jobs=-1,
random_state=1) = 88.7

The Accuracy of our initial model is 88.68% which seems to be decent, but the catch is accuracy is not the best score when you have imbalanced dataset. Since ours is an imbalanced dataset accuracy can be sometime misleading so we will check some more evaluation metric to see whether we have a good model or not.

```
[ ]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.99	0.93	71438
1	0.92	0.47	0.62	17561
accuracy			0.89	88999
macro avg	0.90	0.73	0.78	88999
weighted avg	0.89	0.89	0.87	88999

```
[ ]: confusion_matrix(y_test,y_pred)
```

```
[ ]: array([[70760,    678],
       [ 9381,  8180]])
```

Observations from classification report and confusion matrix:

- The precision is high for class 0.
- The precision of class 1 (defaulter) is also high which is a good thing, high precision for class 1 means if our model predicts that the customer is a defaulter then there is a 92% chance that the person is a defaulter, this might save the bank from defaulters.
- The recall for class 0 is also very high that means it classifies most of the non defaulters correctly.
- The problematic thing is the recall of class 1, basically recall ($tp/(tp+fn)$) tells us how well it can catch the defaulters, since recall for class 1 is very low which implies false negative is high, this means though the person is a defaulter our model is classifying him/her into non defaulter, this might be because of 2 reasons:
 - a) The data is highly imbalanced and there are more records class 0
 - b) We have taken a default threshold of 0.5 currently to classify into class 1 or class 0, which should be changed according to problem statement

Evaluation ROC Curve An ROC curve (receiver operating characteristic curve) is a graph showing the performance of a classification model at all classification thresholds. This curve plots two parameters:

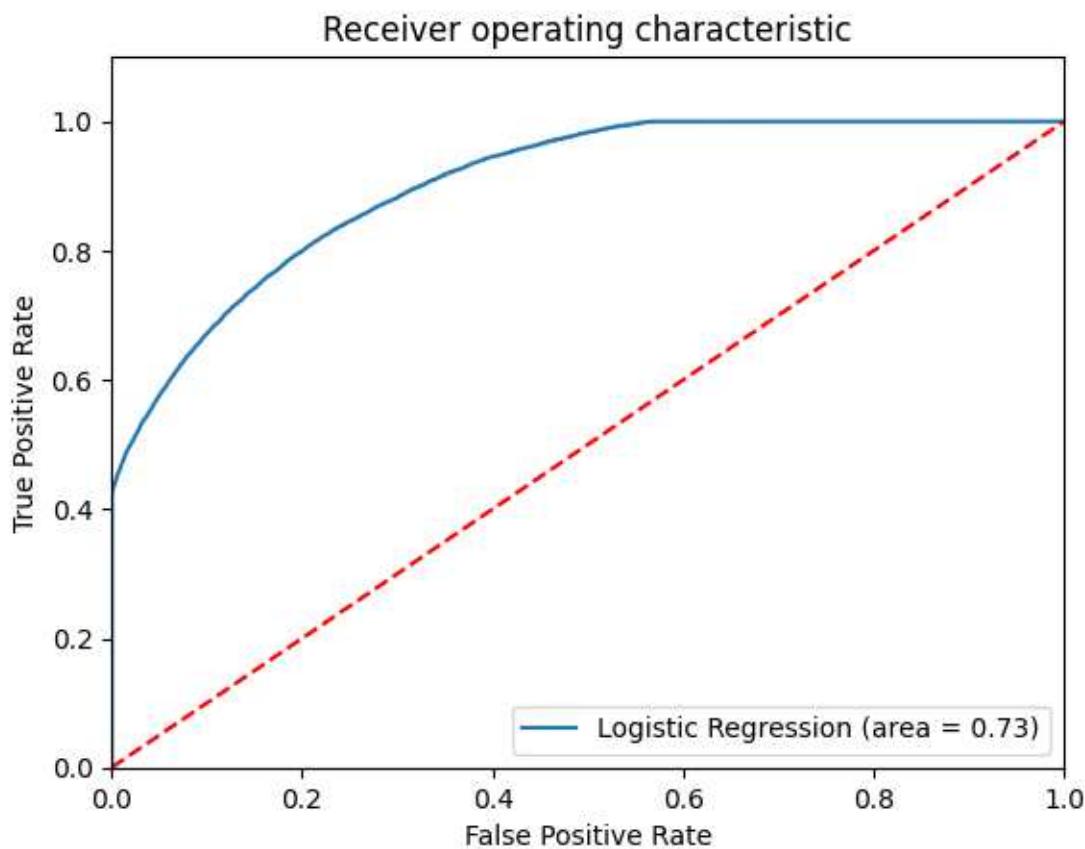
True Positive Rate False Positive Rate True Positive Rate (TPR) is a synonym for recall and is therefore defined as follows:

TPR=(TP)/(TP+FN) False Positive Rate (FPR) is defined as follows:

FPR=(FP)/(FP+TN) An ROC curve plots TPR vs. FPR at different classification thresholds. Lowering the classification threshold classifies more items as positive, thus increasing both False Positives and True Positives. The following figure shows a typical ROC curve.

```
[ ]: roc_auc_score(y_test,y_pred)
fpr,tpr,thr = roc_curve(y_test,model.predict_proba(x_test)[:,1])
```

```
[ ]: score = roc_auc_score(y_test,y_pred)
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % score)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.10])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



The above plot shows us the true positive rate and false positive rates for different thresholds of probabilities.

The higher the area of blue curve the better is our model.

The red dotted line shows a random model which has an `roc_auc_score` of 0.5.

The threshold is chosen based on optimum value of tpr and fpr according to the business problem being solved.

Precision Recall Curve

```
[ ]: precision,recall,thr = precision_recall_curve(y_test,model.  
         ↪predict_proba(x_test)[:,1])  
  
[ ]: # Precision and recall will have one extra value.  
# The extra value in the precision and recall arrays ensures that all relevant  
# decision thresholds are covered, including the boundary cases where the  
# number of predicted positive samples is either zero or equal to the total  
# number of samples.  
len(thr)
```

```
[ ]: 88999
```

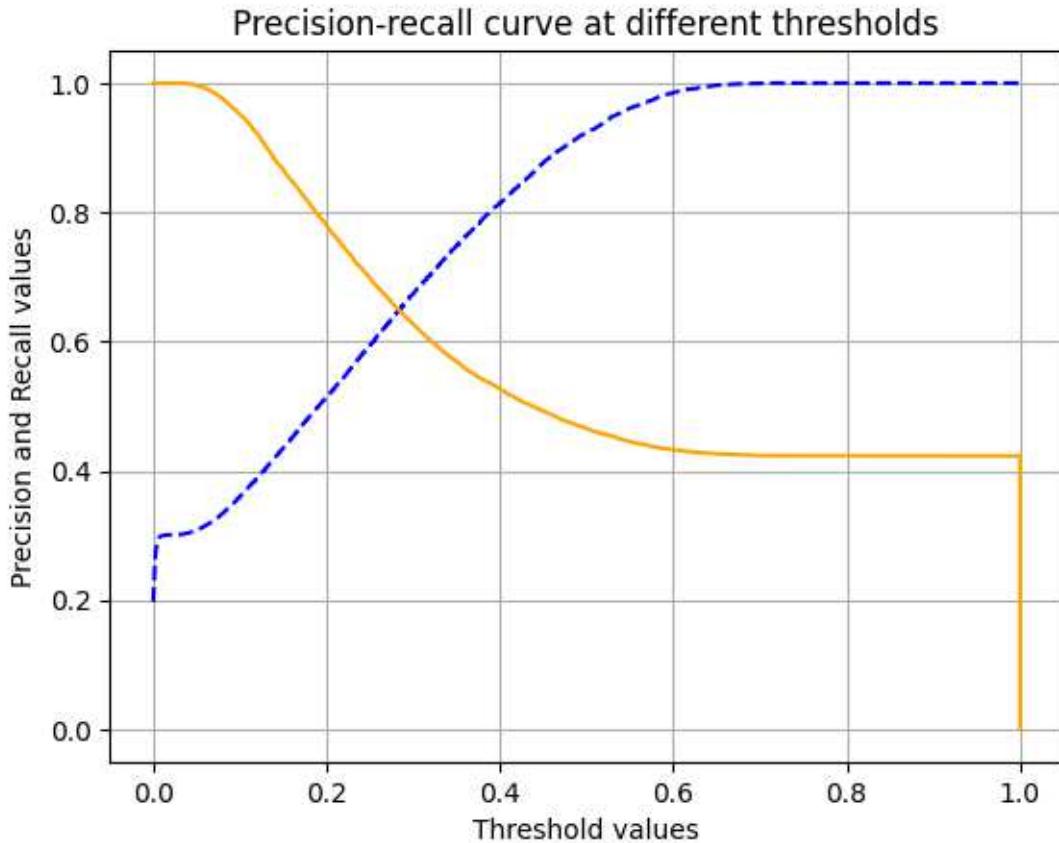
```
[ ]: len(precision)
```

```
[ ]: 89000
```

```
[ ]: len(recall)
```

```
[ ]: 89000
```

```
[ ]: plt.figure()  
plt.plot(thr,precision[0:  
         ↪len(thr)],linestyle='--',label='precision',color='blue')  
plt.plot(thr,recall[0:len(thr)],label='recall',color='orange')  
plt.title("Precision-recall curve at different thresholds")  
plt.xlabel("Threshold values")  
plt.ylabel("Precision and Recall values")  
plt.grid()  
plt.show()
```



The above curve shows us precision and recall values for different thresholds, since both precision and recall are important in our business problem we have to consider a value of 0.3 for the threshold, this means if a model predicts the probability for any person to be more than 0.3 then he/she should be classified as defaulter.

```
[ ]: y_pred = model.predict_proba(x_test)[:,1]

threshold_considered = 0.3

y_pred_custom = (y_pred>threshold_considered).astype('int')
y_pred_custom
```

```
[ ]: array([0, 1, 1, ..., 0, 0, 0])
```

```
[ ]: print(classification_report(y_test,y_pred_custom))
```

	precision	recall	f1-score	support
0	0.91	0.93	0.92	71438
1	0.67	0.63	0.65	17561

```

accuracy           0.87      88999
macro avg         0.79      0.78      88999
weighted avg      0.86      0.87      0.86      88999

```

```
[ ]: confusion_matrix(y_test,y_pred_custom)
```

```
[ ]: array([[66093,  5345],
       [ 6543, 11018]])
```

```
[ ]: # Let's do multicollinearity check and remove some features which have high VIF
    ↵and refit the model
```

Multicollinearity check with VIF

- Multicollinearity occurs when two or more independent variables are highly correlated with one another in a regression model. Multicollinearity can be a problem in a regression model because we would not be able to distinguish between the individual effects of the independent variables on the dependent variable.
- Multicollinearity can be detected via various methods. One such method is Variance Inflation Factor aka VIF. In VIF method, we pick each independent feature and regress it against all of the other independent features. VIF score of an independent variable represents how well the variable is explained by other independent variables.
- $VIF = 1/(1-R^2)$

```
[ ]: from statsmodels.stats.outliers_influence import variance_inflation_factor
def calc_vif(x):
    # Calculating the VIF
    vif = pd.DataFrame()
    vif['Feature'] = x.columns
    vif['VIF'] = [variance_inflation_factor(x.values, i) for i in range(x.
    ↪shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by='VIF', ascending = False)
    return vif

calc_vif(x_train)
```

	Feature	VIF
17	application_type	252.75
4	sub_grade	173.54
3	grade	103.68
6	home_ownership	83.07
2	int_rate	79.69
9	purpose	76.10
8	verification_status	36.23
21	issue_year	31.36
23	earliest_cr_year	19.20

```

15          total_acc    11.44
11          open_acc     11.37
14          revol_util    9.21
7           annual_inc   9.15
10          dti            8.11
0            loan_amnt    6.58
13          revol_bal     5.96
22          issue_month   3.85
24          earliest_cr_month 3.84
5            emp_length    3.41
16          initial_list_status 3.23
18          mort_acc      2.74
19  pub_rec_bankruptcies 2.48
12          pub_rec        2.47
1            term           1.99
20          zipcode        1.63

```

Let's remove the top most vif feature and re-calculate the vif

```
[ ]: removed_features = []
removed_features.append('application_type')
x_train.drop(columns=['application_type'], inplace=True)
```

```
[ ]: calc_vif(x_train)
```

	Feature	VIF
4	sub_grade	172.45
3	grade	103.58
2	int_rate	79.55
9	purpose	57.84
6	home_ownership	54.49
8	verification_status	33.45
20	issue_year	28.62
22	earliest_cr_year	17.91
15	total_acc	11.43
11	open_acc	11.31
14	revol_util	9.06
7	annual_inc	9.00
10	dti	8.10
0	loan_amnt	6.56
13	revol_bal	5.96
23	earliest_cr_month	3.79
21	issue_month	3.75
5	emp_length	3.34
16	initial_list_status	3.09
17	mort_acc	2.57
18	pub_rec_bankruptcies	2.48

```

12          pub_rec    2.47
1           term     1.99
19          zipcode   1.63

```

Since the VIF is very high we can remove 3-4 features at once and can check for VIF again

```
[ ]: removed_features.extend(['sub_grade', 'grade', 'int_rate', 'purpose'])
x_train.drop(columns=['sub_grade', 'grade', 'int_rate', 'purpose'], inplace=True)
```

```
[ ]: calc_vif(x_train)
```

```
[ ]:          Feature      VIF
3       home_ownership 43.12
5   verification_status 31.69
16        issue_year   26.61
18  earliest_cr_year  17.53
11        total_acc   11.40
7         open_acc    11.19
4        annual_inc   8.84
6            dti      8.05
10        revol_util  7.95
0         loan_amnt   6.52
9         revol_bal   5.76
19  earliest_cr_month  3.77
17        issue_month  3.70
2        emp_length   3.30
12  initial_list_status 3.00
13        mort_acc    2.51
14  pub_rec_bankruptcies 2.48
8         pub_rec    2.46
1           term     1.62
15          zipcode   1.61
```

```
[ ]: removed_features.
  ↪extend(['home_ownership', 'earliest_cr_year', 'verification_status', 'issue_year'])
x_train.
  ↪drop(columns=['home_ownership', 'earliest_cr_year', 'verification_status', 'issue_year'], inplace=True)
```

```
[ ]: calc_vif(x_train)
```

```
[ ]:          Feature      VIF
9        total_acc   11.09
5        open_acc    10.64
3        annual_inc   7.87
4            dti      6.99
8        revol_util   6.73
0         loan_amnt   6.14
```

```

7          revol_bal    5.31
15      earliest_cr_month  3.45
14          issue_month   3.29
2          emp_length    3.16
12  pub_rec_bankruptcies 2.47
6          pub_rec       2.44
10  initial_list_status 2.40
11          mort_acc     2.31
1          term         1.60
13          zipcode      1.59

```

```
[ ]: # Now with these features let's refit the model
removed_features
x_test.drop(columns=removed_features,inplace=True)
```

```
[ ]: x_test_cv.drop(columns=removed_features,inplace=True)
```

```
[ ]: model2 = LogisticRegression(max_iter=1000,solver='liblinear',random_state=1,n_jobs=-1)
```

```
[ ]: model2.fit(x_train,y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:1211:
UserWarning: 'n_jobs' > 1 does not have any effect when 'solver' is set to
'liblinear'. Got 'n_jobs' = 2.
warnings.warn(
```

```
[ ]: LogisticRegression(max_iter=1000, n_jobs=-1, random_state=1, solver='liblinear')
```

```
[ ]: print(classification_report(y_test,model2.predict(x_test)))
```

	precision	recall	f1-score	support
0	0.88	1.00	0.93	71438
1	0.97	0.44	0.60	17561
accuracy			0.89	88999
macro avg	0.93	0.72	0.77	88999
weighted avg	0.90	0.89	0.87	88999

```
[ ]: confusion_matrix(y_test,model2.predict(x_test))
```

```
[ ]: array([[71220,    218],
 [ 9862,  7699]])
```

Oversampling using SMOTE to solve the imbalance issue

```
[ ]: from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state=1,n_jobs=-1,k_neighbors=5)

[ ]: x_train_new,y_train_new = sm.fit_resample(x_train,y_train)

/usr/local/lib/python3.10/dist-
packages/imblearn/over_sampling/_smote/base.py:336: FutureWarning: The parameter
`n_jobs` has been deprecated in 0.10 and will be removed in 0.12. You can pass
an nearest neighbors estimator where `n_jobs` is already set instead.
warnings.warn(
[ ]: print(f"After OverSampling, the shape of train_X:{x_train_new.shape}")
print('After OverSampling, the shape of train_y: {}'.format(y_train_new.
    ↪shape))

print(f"After OverSampling, counts of label '1': {sum(y_train_new == 1)}")
print(f"After OverSampling, counts of label '0':{np.sum(y_train_new==0)}")
```

After OverSampling, the shape of train_X:(321892, 16)
After OverSampling, the shape of train_y: (321892,)

After OverSampling, counts of label '1': 160946
After OverSampling, counts of label '0':160946

```
[ ]: model3 = LogisticRegression(max_iter=1000,random_state=1,n_jobs=-1,)
model3.fit(x_train_new, y_train_new)
predictions = model3.predict(x_test)

# Classification Report
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.93	0.80	0.86	71438
1	0.48	0.77	0.59	17561
accuracy			0.79	88999
macro avg	0.71	0.78	0.73	88999
weighted avg	0.84	0.79	0.81	88999

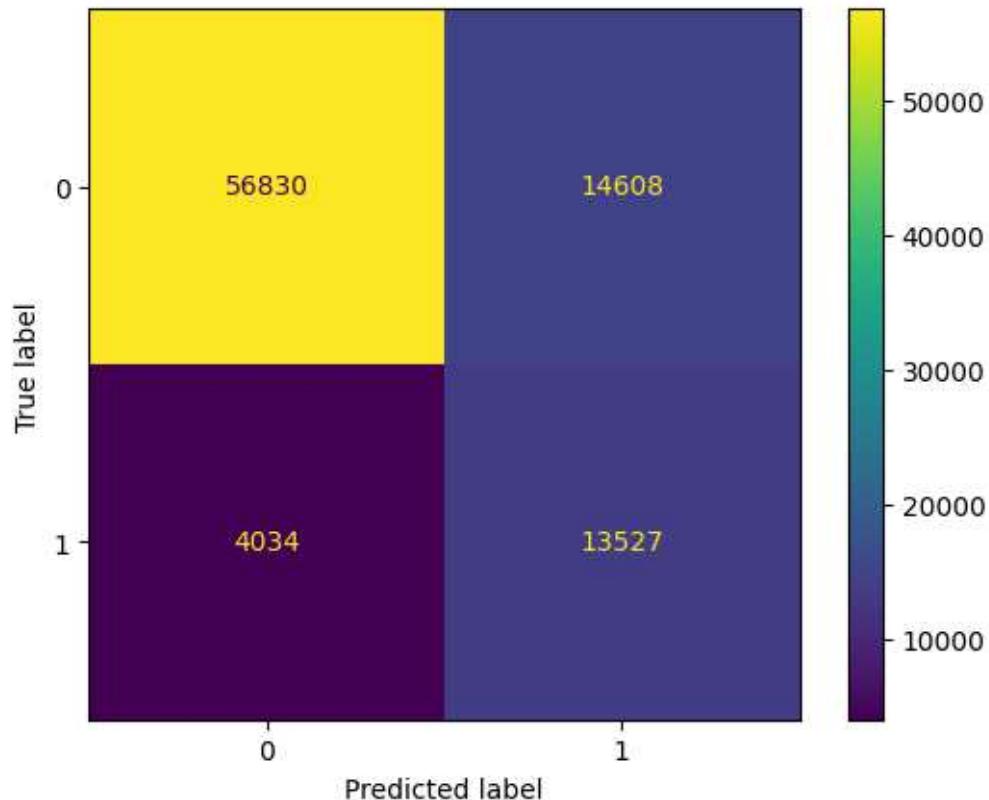
```
[ ]: cm=confusion_matrix(y_test,predictions)

[ ]: model3.score(x_test,y_test)

[ ]: 0.7905369723255318
```

```
[ ]: from sklearn.metrics import ConfusionMatrixDisplay  
ConfusionMatrixDisplay(cm).plot()
```

```
[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at  
0x7b8ba3e626e0>
```



0.1 Insights and Recommendations

Insights

1. Loan Grade Concentration: The concentration of loans in the B, C, and A categories suggests these grades are most common and potentially the most reliable segments.
2. Employment Stability and Loan Purpose: The long employment duration of many borrowers indicates stability, which is a positive sign for lenders. The primary purpose of loans being debt consolidation reflects a trend towards financial management and restructuring by borrowers.
3. Correlation Insights: The high correlation between loan amount and installment (0.95 Pearson, 0.97 Spearman) suggests that these two variables are strongly linked in their behavior. The minimal difference in the pearson and spearman correlation coefficients indicate that there is minimal to no involvement of non linearity in the data due to which did not check for any polynomial features.

4. Verification Status: The fact that verification doesn't necessarily guarantee a loan being fully paid implies that other factors play a significant role in loan performance.

Recommendations

1. Tailored Loan Products: Develop tailored loan products for the most common borrower segments (B, C, and A grade borrowers) to enhance product-market fit.
2. Loan Term Structuring: Offer more flexible terms for higher-grade loans and consider stricter terms for longer-term and lower-grade loans to mitigate risk.
3. Further Statistical Analysis: Conduct further statistical tests to validate the significance of observed correlations and insights, ensuring that lending strategies are data-driven.
4. Monitoring and Adjustment of Models: Continuously monitor and adjust credit scoring models in response to changes in borrower behavior and economic conditions.